

Task 1 Report by Ryan Wickman

Experiment Settings

Binary Classification Hyperparameters

Hyperparameter	Value
Optimizer	Adam
Learning Rate	0.01
LR Decay Rate	0.99
Epochs	100
Batch Size	16
Early Stopping	True
Patience	10
Num. Random init	3
K-fold Split Size	5

Multi-class Classification Hyperparameters

Hyperparameter	Value
Optimizer	Adam
Learning Rate	0.01
LR Decay Rate	0.99
Epochs	100
Batch Size	256
Early Stopping	True
Patience	10
Num. Random init	3
K-fold Split Size	N/a

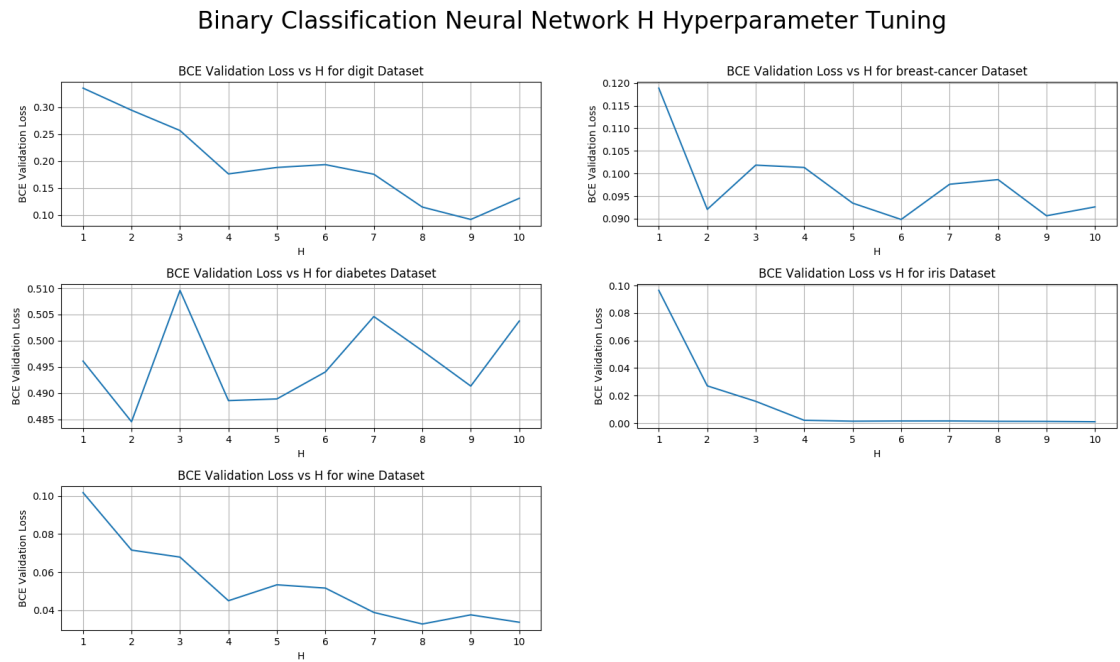
Parameter Tuning Results

Binary Classification Tuning Results

I used k-fold cross-validation to evaluate the number of units in the hidden layer, H , that provides the lowest validation loss. In my experiments, H was tested on values in the range of $[1,10]$.

The table in the Binary Classification Hyperparameters section above provides details on how the experiments were setup.

The figure below displays the BCE Validation Loss vs H plot for every dataset:



From the graph above, you can easily evaluate the best H value for every dataset.


Multi-class Classification Tuning Results

I used normal cross-validation, a simple 80/20 split between the training and validation set. I used this to evaluate the number of units in to use in first hidden layer and the second hidden layer, L1 and L2; respectively.

The table in the Multi-class Classification Hyperparameters section above provides details on how the experiments were setup.

The figure below displays both the CE Validation Loss vs (L1, L2) plot and the CE Validation Loss vs (L1, L2) plot, where every combination of L1 and L2 is included:

Multi-class Classification Neural Network (L1, L2) Hyperparameter Tuning



From the graph above, you can easily evaluate the best L1 and L2 combination. When L1 = 75 and L2 = 20, we get the minimum loss and the maximum accuracy.